

PERFORMANCE ENGINEERING

Lecture 6: Statistical models and Performance counters

May 9th, 2022

Ana Lucia Varbanescu
a.l.varbanescu@uva.nl
(room C3.109)

Today

- Data-centric/statistical models
 - ... and examples

Remaining topics

- Bottleneck analysis (L7)
- Large scale systems modelling and analysis (L7)
- Queuing theory and its applications to PE (L8)
- Simulators (L8/L9)
- The polyhedral model and its applications to PE (L8/L9)

No lecture/lab in the last week (May 29th), focus on project and exam.

Reminder ...

- **Analytical modeling**

- **Goals:** “can I predict the performance of my system/application?”
 - System = application + data + hardware
- **Idea:** model the operation of the system in a symbolic model and calibrate for real system

- **Data-centric/Statistical modeling**

- **Goals:** “can I predict the performance of my application for an unseen input and/or unseen hardware?”
- **Idea:** use past execution to build a model of the application performance

- **Simulation-based modeling (next week)**

- **Goals:** “can I predict the performance of my system?”
 - System = application + data + hardware
- **Idea:** simulate the operation of the system and measure different events

Statistical/data-centric/data-driven
performance modeling

What's in a name?

- Data-driven/data-centric/statistical/machine-learning/... modeling: building a model based on historical data to predict future behavior.
 - Typically considered some form of *black-box modeling*.
- Requirements:
 - Historical/representative data
 - Statistical/machine-learning approaches
 - (Linear) regression
 - Decision tree(s)
 - Random forest
 - Neural networks
 - ...

Why (not) statistical modeling?

Pro:

- Analytical modeling cannot capture all the complexities of the system
 - Data-dependent behavior
 - Hardware-specific behavior
- Faster than analytical modeling and simulation
 - Modulo training time ...
- Could provide additional information/intuition on the inner workings of the system

Con:

- Sensitive to incorrect/insufficient data
- Sensitive to the statistical method
- Not automatically insightful

Basic idea / workflow

1. Determine **target variable**.
 - What to predict (some performance metric).
 2. Collect **sufficient representative** performance data points
 - Obtain a labeled training set.
 3. Determine **features (aka, predictor variables)**.
 - What does the performance depend on.
 4. Choose a **modeling technique** with enough accuracy
 5. Train your model
 - Random-split or cross-validation
 6. Validate your model
- Challenges?

1. Target variable

- Any performance metric ...
 - Supervised learning => labeled data => must have it as part of the training set!
- Examples
 - Execution time
 - Cycles
 - Instructions per cycle
 - Ranking of performance
 - Atomic contention
 - Cache miss/hit
 - ...

2. Obtain/collect labeled training data

- Information about past performance
- Collection requires
 - Instrumentation
 - Execution
 - Statistics on feature and target variable
 - ... and proper storage!
- How much data?
 - Depends on the method you choose/aim for
- Representative?
 - Difficult to assess ...

3. Selecting features

- Naïve feature selection
 - Select all parameters that might impact performance
- Too many? Feature pruning!
 - Sensitivity analysis
 - Check with specific datasets whether some features are indeed useful
 - Variable importance (as part of different machine learning models)
 - Combined features
- Too few?
 - Collect more (detailed) statistics.

4. Choose a modeling technique

- Regression – most commonly used
 - Simple linear regression
 - General polynomial regression
- Decision trees
- Random forest
- ...
- How to select?
 - Use the simplest that match the expected behavior
 - Use the one that needs the least data
 - Use the one that is less sensitive to overfitting
 - Use the one that provides added benefits (e.g., variable importance)
 - Use the most generic one
 - Use the fastest to apply

5. Train the model

- As simple as running the model training functions on the test data
 - Typically an iterative process
 - Implications on steps 4 (technique) and 3 (features)
- Be careful
 - Make sure prediction variables match method requirements
 - Use good practices for the selection of test and training data
 - Revisit predictor variables if needed
 - Store model and outcome analysis if possible
 - Model itself for post-analysis
 - Variable importance

6. Validate the model


- Apply the model on the test set.
- Report/analyze errors.
- Revisit 3,4,5 for tweaking.
- Revisit 1,2 in case accuracy is really low.
- Save final model for further use!

Example 1: Diffusion Monte Carlo (DMC)

```
For generation = 1 ... Nsteps Do
  For walker = 1 ... Nwalk Do
     $\mathbf{r} := \{\mathbf{r}_1, \dots, \mathbf{r}_n\}$ 
    For electron  $i = 1 \dots n$  Do
      Set  $\mathbf{r}'_i = \mathbf{r}_i + \delta$ 
      Compute  $D(\mathbf{r}') = D(\mathbf{r}_1, \dots, \mathbf{r}'_i, \dots, \mathbf{r}_n)$ 
      If Accepted Then
        Update Invers  $B(\mathbf{r}) \leftarrow B(\mathbf{r}')$ 
      End If
    End Do
    Compute Energy
  End Do
  Branch and Accumulate Averages
End Do
```


DMC: $T = f(\text{steps})$

Execution time (w=8192)



DMC: $T = f(\text{walkers})$

Execution time (s=150)




DMC: conclusions

- Execution time scales linearly with the number of steps and the number of walkers.
- Minor errors at small simulation sizes
 - To be investigated in case small sizes are relevant
- Lessons learned:
 - Linear regression was sufficient
 - Target (execution time) and predictor variables (input size) obvious

Example 2: BFS


- Graph traversal algorithm
- Visits nodes in levels
- Many parallel algorithms exist
 - Some are platform specific!
- Performance is known to be graph dependent
 - Topology impacts execution time
 - Start node impacts execution time
- Challenges?
 - How to measure topology
 - How to get sufficient representative training data



BFS: experimental setup

- GPU platforms
- 15 different versions of BFS
 - Some are variants, some are truly different
 - Differ in amount of parallelism (vertex-based vs. edge-based)
 - Differ in synchronization mechanisms (atomics vs. lock-free)
- ~200 graphs from KONECT
 - Different starting nodes
- Collected statistics
 - Execution time (total)
 - Execution time (per level)
 - Level & next level size
 - *Performance counters* (added later)*

Normalized performance



BFS: attempt 0

- Analytical modeling
- How would you build that?

BFS: attempt 1

- Target variable: execution time per algorithm
- Feature variables: graph properties
 - Number of edges
 - Number of vertices
 - Diameter
 - ...
- Random forest => failure
- Decision tree => failure
- (many other failures)

What can we do?

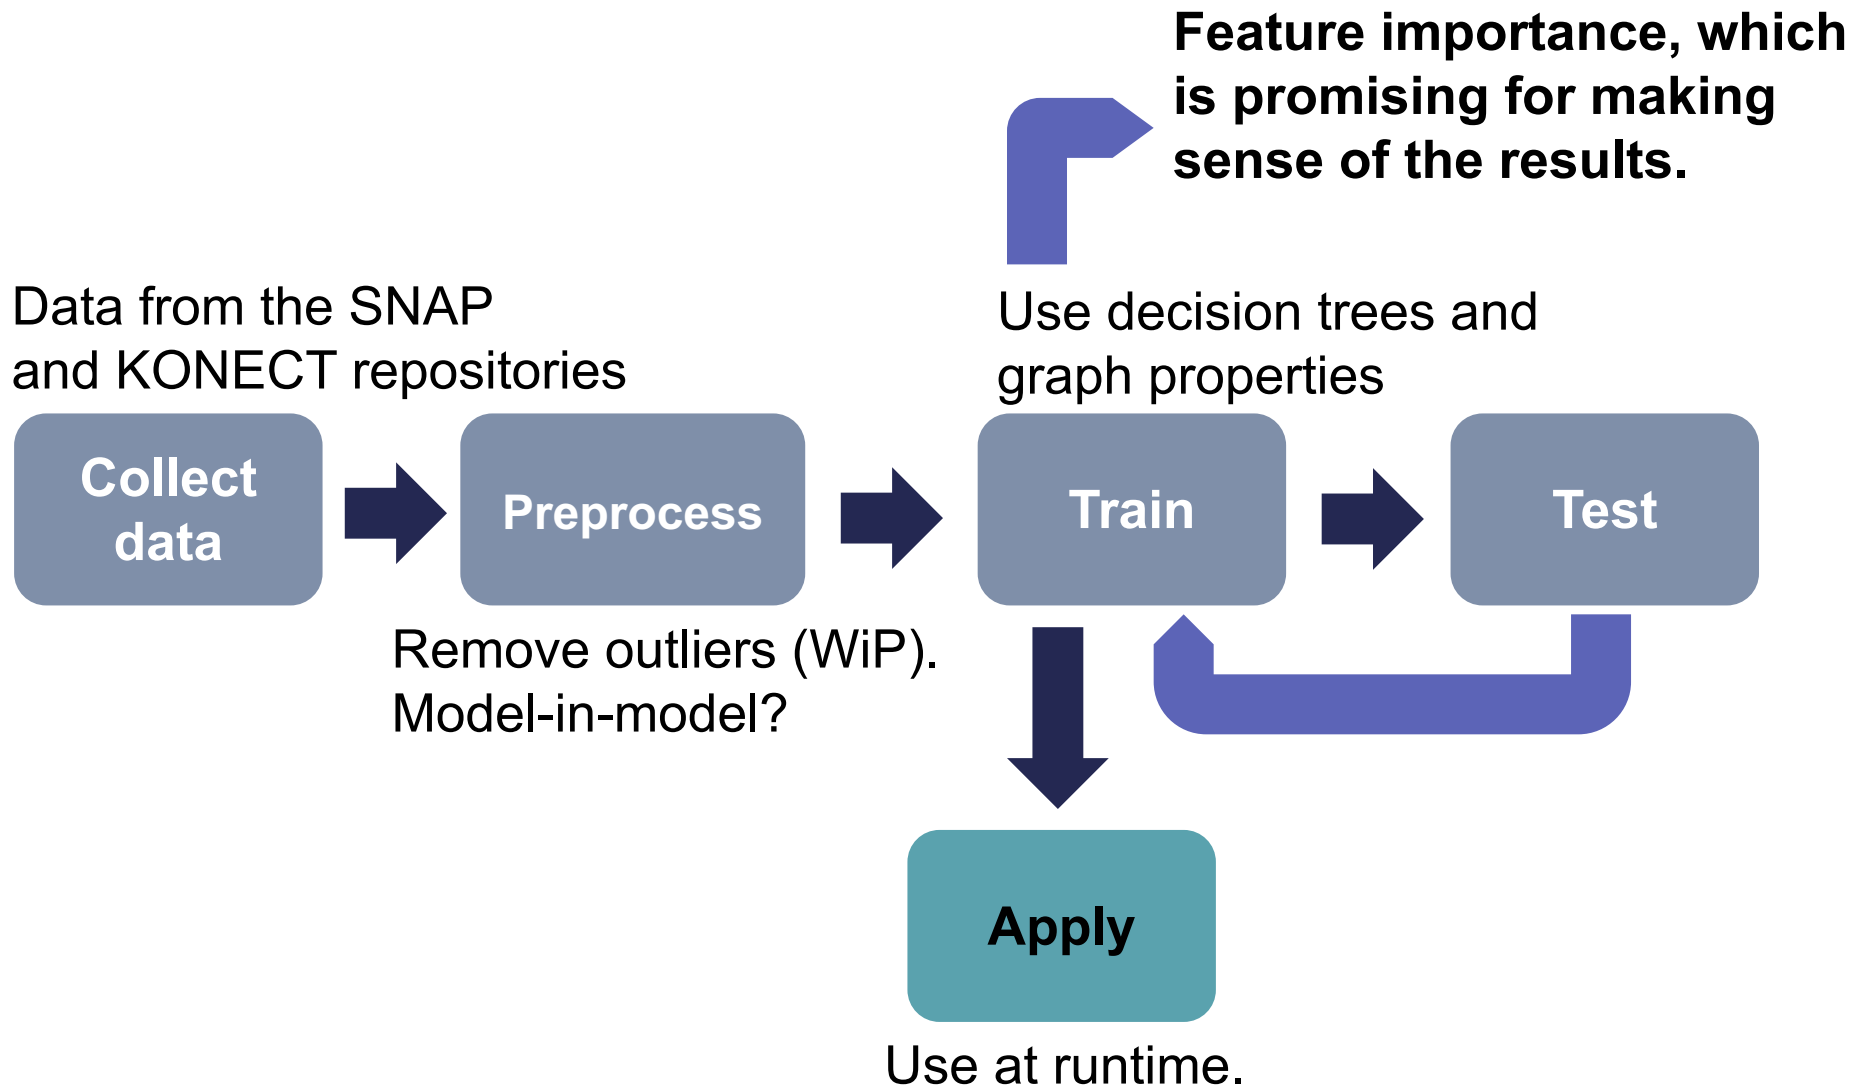
BFS: attempt 2

- Target variable: execution time per level
- Feature variables: graph properties
 - Number of edges
 - Number of vertices
 - Level and frontier size
 - Percentage of graph visited
- Analytical model: failure due to atomics impact
- Regression: not accurate enough
- Random forest:
 - High accuracy
 - High prediction cost

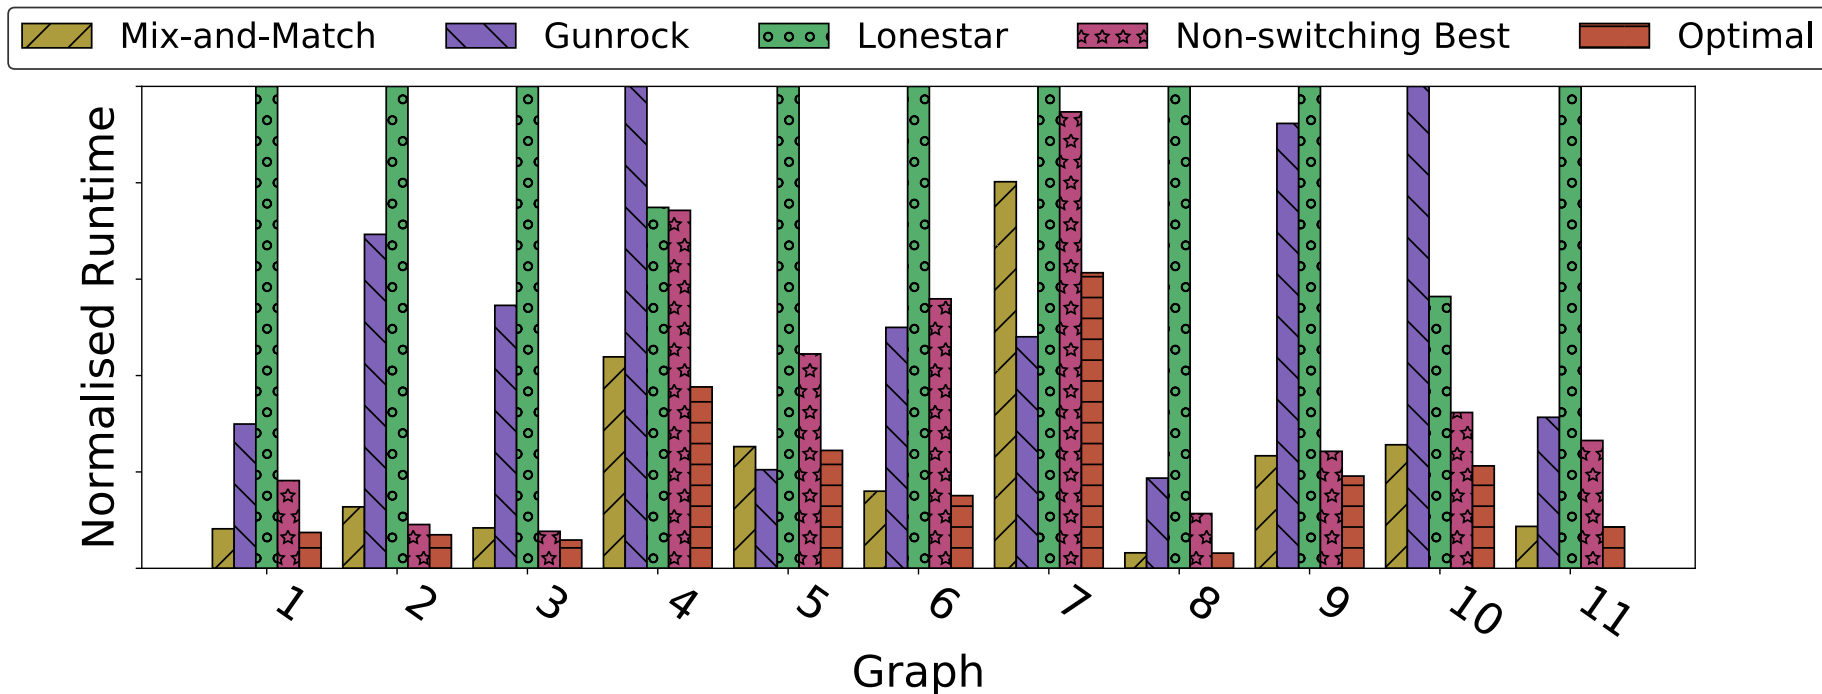
BFS: attempt 3

- Target variable: ranking of best algorithm
- Feature variables: graph properties
 - Number of edges
 - Number of vertices
 - Level and frontier size
 - Percentage of graph visited
- Random forest:
 - High accuracy
 - High prediction cost
- Decision tree: worked!

BFS: Current workflow



Does it really work?

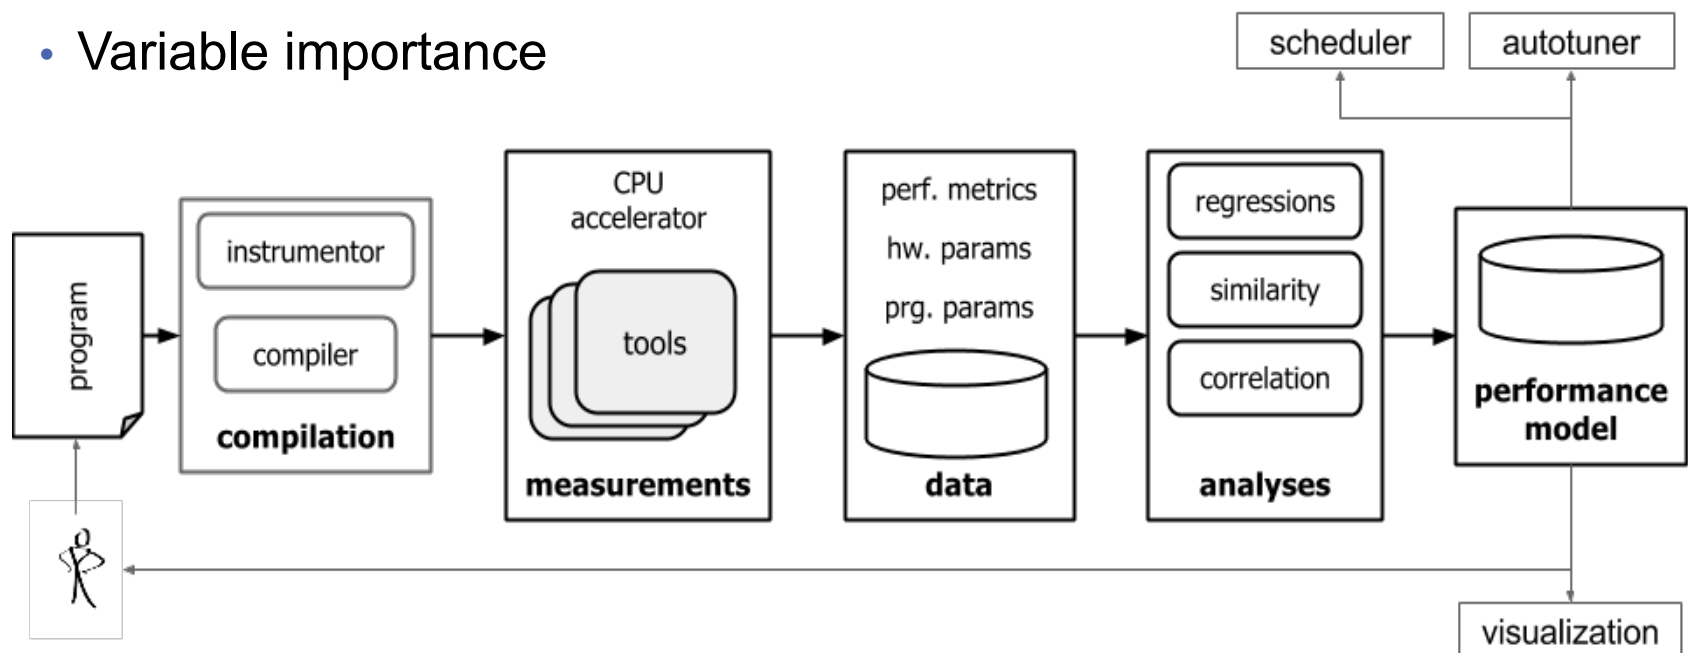


- Runtime switching is possible, (currently) with some memory overhead
- We are faster than the state-of-the art, on average, by 3x

Mix-and-match uses performance variability to build the best BFS per graph!

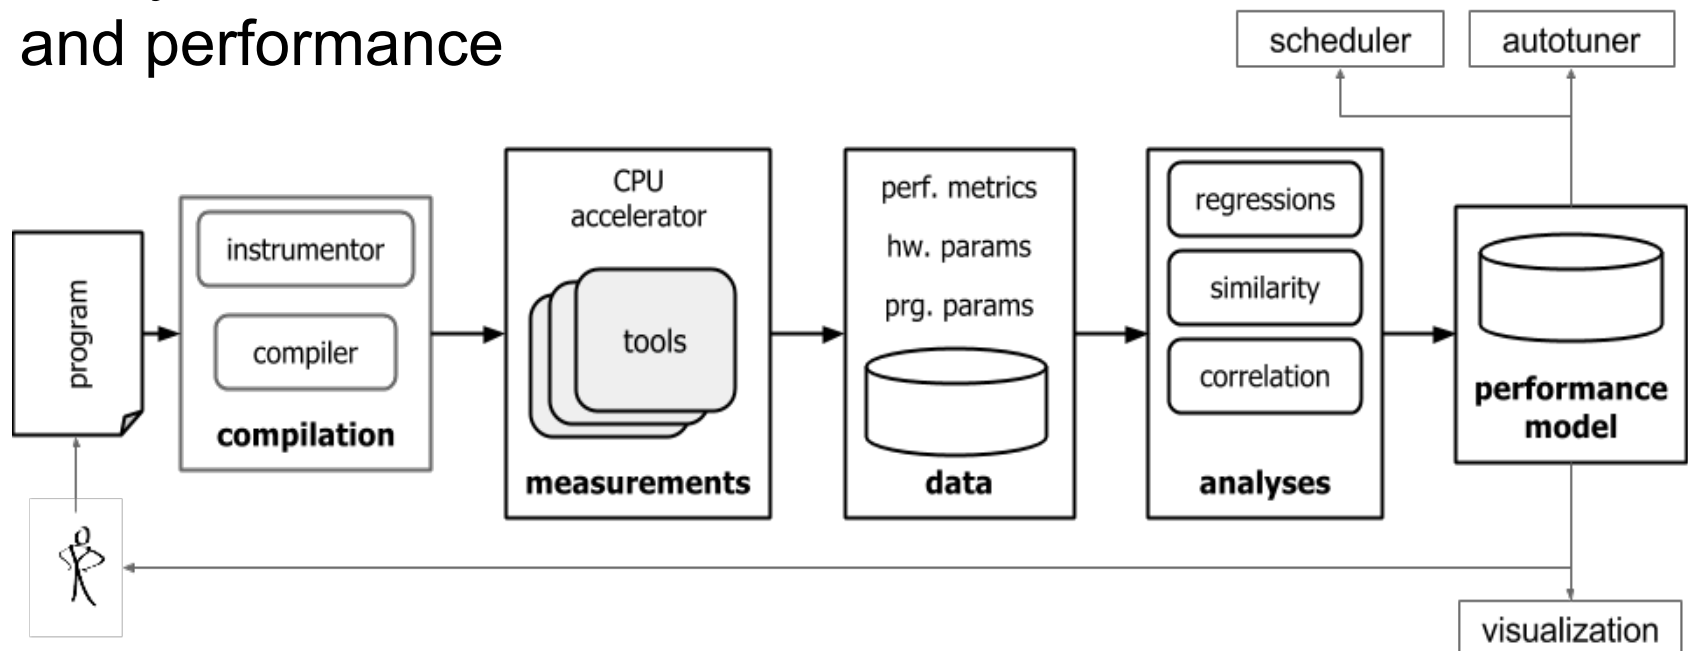
Example 3: the BlackForest framework

- Automate the process of statistical modeling
- Target variable: execution time
- Predictor variables: performance counters
- Outcome:
 - Prediction model
 - Variable importance




Example 3: the BlackForest Framework

- Compilation: optional, scope limitation by instrumentation
- Measurements: performance data collection via hardware performance counters
- Data: repository, file system, database
- Analyses: reveal correlation between counter behavior and performance

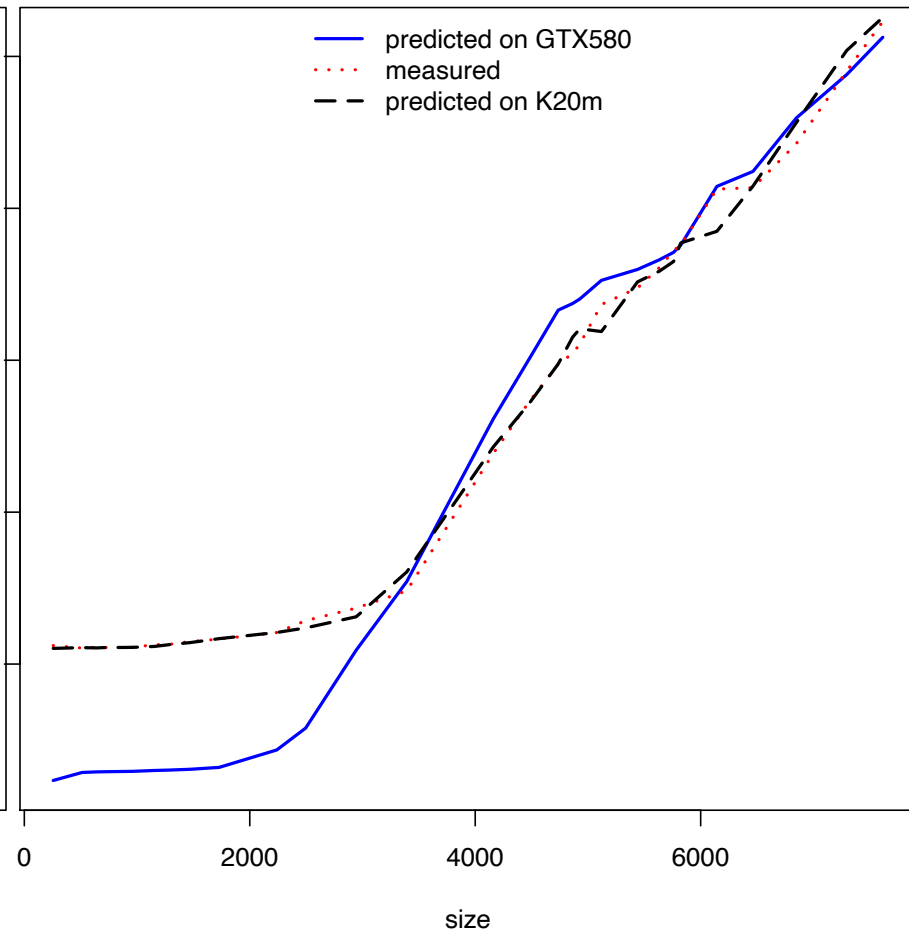


BlackForest: Results*


Matrix multiply



Needleman-Wunsch



BlackForest: Variable importance



Statistical Modeling recap

- Statistical approaches
 - Black-box modeling
 - Useful when “opening” the system is not feasible
 - Time or complexity issues
- Requirements
 - Labeled training data
 - Tools for training and building models
- Challenges
 - Getting/collecting representative data
 - Variable selection
 - Validation & testing
- Disadvantages
 - Low insight
 - Expensive re-training



More data for your model ...

What if ...

- ... the data is too coarse-grain?
- ... we can't understand the dependency of the target variable on the input data?
- ... we can't properly calibrate?
- ... we have data-dependent behavior?

We can get more insight into the performance behavior/causes for it using more detailed analysis of past execution by **monitoring hardware events**.