

PERFORMANCE ENGINEERING

Lecture 8: Large-scale systems and applications

May 19th, 2022

Ana Lucia Varbanescu
a.l.varbanescu@uva.nl

Quizzes last time(s) – L7Q1

Consider SpMV.

Assume the matrix is 4096 x 4096. You split the matrix over the 16 cores on a DAS node, row-wise. The vector is replicated.

You then transform the matrix in CSR, perform computation, and concatenate results.

- What is the main performance problem?
- What PC's you need for it?

L7Q1 – Performance counters

- Main problem:
 - **Load imbalance**
 - NUMA problems
 - Cache (false) sharing (?)
- Identify it:
 - **Check counter on instructions (executed/retired)**
 - Check sleep/idle status
 - Check remote data transfers

Quizzes last times – L6Q1

- Please design a statistical model to estimate the performance of a parallel histogram for a single (very large) image.
- Please explain how you will design your model, considering the following questions.
 - What output?
 - What features?
 - Collect training data?
 - Validation? Accuracy?

L6Q1

- Output
 - Execution time / number of cycles for a given image
 - Some form of contention / delay due to false sharing or other issues
- Input (features)
 - dimensions of the image (width and height)
 - distribution of the pixel colors (if atomic operations are used).
 - Image entropy
 - number of threads
 - machine features
- Collect data
 - Images (generated/downloaded)
 - ... and their timing
 - Cost?
- Validation
 - Compare to measurement and compute error for unseen images

Today

- What about large-scale systems and applications?
- Models, tools, examples ...
- ... and demo.

Distributed (memory) systems


Distributed (memory) systems

- Same principle for performance modeling:

$$T = F (T_compute, T_comm, T_sync) =$$

$$F (f(T_compute_node(i)) , g(T_comm(i)) , h(sync))$$

$$F = ? , f,g,h = ?$$



Distributed (memory) systems

- Same principle for performance modeling:

$$T = F (T_{\text{compute}}, T_{\text{comm}}, T_{\text{sync}}) = \\ F (f(T_{\text{compute_node}(i)}) , g(T_{\text{comm}(i)}) , h(\text{sync}))$$


$F = +$ (BSP model)

$T = f (T_{\text{compute_node}(i)}) + g() + h()$

$F = \max$ (overlapping compute and communicate)


$T = \max (f(), g(), h())$

- Boils down to:
 - Understand synchronization model
 - Understand communication model



Typical programming

- Message passing (e.g., MPI)
- In message passing, all communication is explicit!
 - Advantage?
 - Disadvantage?




Cost of communication

- Message size
 - The larger, the more time consuming


- Message type
 - Point-to-point
 - Collective communication

- Network topology


- Not enough knowledge?




broadcast



scatter



gather



reduction



Ring



Mesh




Star



Fully Connected



Line



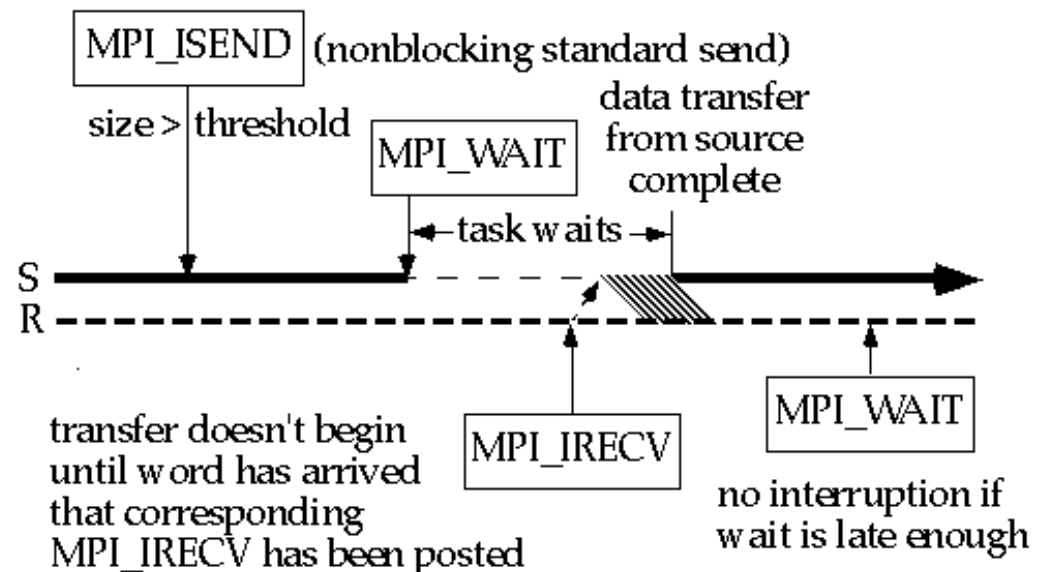
Tree



Bus


Cost of synchronization

- Blocking vs. non-blocking messages
 - Who is blocking ?
 - Who pays the overhead
 - When is it unblocking ?
 - How much overlap happens?
 - What is the cost of a wait?



Cost of synchronization

- Barriers
 - Global
 - Per groups



The good news ...

- Models focus on scaling behaviour and extrapolation
 - And its potential bottlenecks
- Most current models are a mix of ...
 - High-level analytical models
 - Communication pattern analysis
 - Heavy calibration
- There is increasing tool support for such applications.
- See presentation by Dr. Alexandru Calotoiu, ETH Zurich
 - Here:
https://canvas.uva.nl/files/6728556/download?download_frd=1